

finalbzhangfinalfinal

JIASHU MIAO

6/4/2019

```
datafull <- read.csv("full.csv")
# full <- datafull
dataz2 <- datafull
# datafull$Gender %>% table()
datafull$Gender <- ifelse(datafull$Gender=="Female",1,0)
# datafull$Gender

pei=data.frame(table(datafull$READ))
pei

## [1] Freq
## <0 rows> (or 0-length row.names)

7331    /68890

## [1] 0.106416

library(tidyverse) # data manipulation and visualization

## — Attaching packages — tidyverse 1.2.1 —

## [1] ggplot2 3.1.0      [1] purrr  0.2.5
## [1] tibble  2.1.1      [1] dplyr  0.7.8
## [1] tidyr   0.8.2      [1] stringr 1.3.1
## [1] readr   1.3.1      [1] forcats 0.3.0

## — Conflicts — tidyverse_conflicts() —
## [1] dplyr::filter() masks stats::filter()
## [1] dplyr::lag()    masks stats::lag()

library(modelr) # provides easy pipeline modeling functions
library(broom, warn.conflicts = F)
library(ggplot2)
library(caret)

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift
```

```

library("pROC")

## Type 'citation("pROC")' for a citation.

##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##      cov, smooth, var

library("ranger")

dataz2$los_days <- as.integer(as.character(dataz2$los_days))

## Warning: NAs introduced by coercion

dataz2$L0S_Score <- dataz2$L0S_Score %>% as.character() %>% as.integer()

## Warning in function_list[[k]](value): NAs introduced by coercion

#check there is NA
dim(dataz2)

## [1] 77000    23

dataz2 <- dataz2[!(is.na(dataz2$L0S_Score)),]
# we take out 11 rows that contains NA and they are just a few so we could remove simply remove the rows

any(colSums(is.na(dataz2))>0) #no NAs now.

## [1] FALSE

SM_inf=0
for (i in 1:23)
{ SM_inf <- sum(is.infinite(dataz2[,i]))
  if (!SM_inf==0)
    print(SM_inf)
}

SM_null=0
for (i in 1:23)
{ SM_null <- sum(is.null(dataz2[,i]))
  if (!SM_null==0)
    print(SM_null)
}

SM_nan=0
for (i in 1:23)
{ SM_nan <- sum(is.nan(dataz2[,i]))
  if (!SM_nan==0)
    print(SM_nan)
}

```

```

}
dim(dataz2)

## [1] 76221    23

datafull=dataz2

# datafull$CMS_Readmission_unplanned_flag %>% Levels()
datafull$READ <- as.integer(datafull$CMS_Readmission_unplanned_flag)-2
datafull$Clarity.LACE..Score <- as.integer(datafull$Clarity.LACE..Score)

set.seed(123)
train.index <- createDataPartition(datafull$READ,p = .75, list = FALSE)
new_train <- datafull[train.index,]
new_test <- datafull[-train.index,]

dim(new_train)

## [1] 57166    24

dim(new_test)

## [1] 19055    24

new_train %>% names()

## [1] "Gender"                                "Age"
## [3] "EDV_Score"                            "cadm_type_score"
## [5] "LOS_Score"                            "Charlson_Score"
## [7] "ALC_Score"                            "Elective_adm_Score"
## [9] "Urgent_adm_Score"                     "Teach_Score"
## [11] "Male_Score"                           "Admit_DT"
## [13] "Discharge_DT"                         "los_days"
## [15] "ADMIT_SOURCE"                         "dis_from_base_class"
## [17] "loc_name"                             "DISCH_DISPOSITION"
## [19] "marital_status"                       "PAT_HOMELESS_YN"
## [21] "zip"                                  "CMS_Readmission_unplanned_flg"
## [23] "Clarity.LACE..Score"                  "READ"

```

```

## single
set.seed(123)
single_logistic_model <-
  glm(
    data = new_train,
    formula = READ~Clarity.LACE..Score,
    binomial(link="logit"))
summary(single_logistic_model)

##
## Call:
## glm(formula = READ ~ Clarity.LACE..Score, family = binomial(link = "
logit"),
##     data = new_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.7756  -0.5293  -0.3963  -0.2947   2.8569
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -4.1316547   0.0550712  -75.02  <2e-16 ***
## Clarity.LACE..Score  0.0338960   0.0008833   38.38  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 36265  on 57165  degrees of freedom
## Residual deviance: 34513  on 57164  degrees of freedom
## AIC: 34517
##
## Number of Fisher Scoring iterations: 5

exp(coef(single_logistic_model))


##           (Intercept) Clarity.LACE..Score
##           0.01605629      1.03447697

single_log_pre <- predict(single_logistic_model,new_test, type= "respon
se")
length(single_log_pre)

## [1] 19055

roccurve <- roc(READ~single_log_pre,data = new_test, plot=TRUE, grid=T
RUE, print.auc=TRUE)


```



```
print(roccurve)

##
## Call:
## roc.formula(formula = READ ~ single_log_pre, data = new_test,      pl
ot = TRUE, grid = TRUE, print.auc = TRUE)
##
## Data: single_log_pre in 17237 controls (READ 0) < 1818 cases (READ
1).
## Area under the curve: 0.6764

plot(roccurve,xlim=c(0,1),col="blue",asp = NA)
```



```
coords(roc=roccurve,"best","threshold")

## threshold specificity sensitivity
## 0.1013310 0.6268492 0.6452145

classifier <- as.numeric(coords(roc=roccurve,"best","threshold"))
# Area under the curve: 0.6764
class <- ifelse(single_log_pre>classifier[1],1,0)
accuracy <- 1-mean(class!=new_test$READ)
cat("Accuracy is for simple logistic ", accuracy,"\n")

## Accuracy is for simple logistic 0.6286014

cat("The AUC score is ",roccurve$auc,"\n")

## The AUC score is 0.6763661

set.seed(16)
#Multivariate
# multi_logistic_model <-
#   glm(
#     data = new_train,
#     formula = READ~Clarity.LACE..Score+Age+Los_days+,
#     family=binomial(link="logit")
```

```

# )

multi_logistic_model2 <-
  glm(
    data = new_train[,c(3:11,24)],
    formula = READ~.,
    family=binomial(link="logit")
  )
summary(multi_logistic_model2)

##
## Call:
## glm(formula = READ ~ ., family = binomial(link = "logit"), data = ne
w_train[,
##      c(3:11, 24)])
##
## Deviance Residuals:
##      Min        1Q    Median        3Q        Max
## -1.4958  -0.4803  -0.3282  -0.2264   3.0136
##
## Coefficients: (2 not defined because of singularities)
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -4.530269    0.057960  -78.162 < 2e-16 ***
## EDV_Score       0.188342    0.006842   27.529 < 2e-16 ***
## cadm_type_score  0.025770    0.003041    8.473 < 2e-16 ***
## LOS_Score       0.177146    0.006051   29.277 < 2e-16 ***
## Charlson_Score  0.070337    0.005865   11.993 < 2e-16 ***
## ALC_Score              NA           NA      NA      NA
## Elective_adm_Score 0.035108    0.007406    4.741 2.13e-06 ***
## Urgent_adm_Score  0.019324    0.001425   13.558 < 2e-16 ***
## Teach_Score              NA           NA      NA      NA
## Male_Score        0.006149    0.009941    0.619  0.536
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 36265  on 57165  degrees of freedom
## Residual deviance: 31944  on 57158  degrees of freedom
## AIC: 31960
##
## Number of Fisher Scoring iterations: 6

exp(coef(multi_logistic_model2))

##      (Intercept)      EDV_Score      cadm_type_score
##      0.01077777      1.20724682      1.02610531
##      LOS_Score      Charlson_Score      ALC_Score
##      1.19380532      1.07286931      NA

```

```

## Elective_adm_Score    Urgent_adm_Score    Teach_Score
##           1.03573184           1.01951167           NA
##           Male_Score
##           1.00616761

multi_log_pre <- predict(multi_logistic_model2,new_test[,c(3:11,24)], type= "response")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

length(multi_log_pre)


## [1] 19055

roccurve_multi <- roc(READ~multi_log_pre,data = new_test)
print(roccurve_multi)

##
## Call:
## roc.formula(formula = READ ~ multi_log_pre, data = new_test)
##
## Data: multi_log_pre in 17237 controls (READ 0) < 1818 cases (READ 1).
## Area under the curve: 0.7588

plot(roccurve_multi,xlim=c(0,1),col="blue",asp = NA)

```

```
coords(roc=roccurve_multi,"best","threshold")

##   threshold specificity sensitivity
## 0.08967944 0.65579857 0.74642464

classifier_multi <- as.numeric(coords(roc=roccurve_multi,"best","threshold"))
# Area under the curve: 0.6764
class_multi <- ifelse(multi_log_pre>classifier_multi[1],1,0)
table(class_multi)

## class_multi
##      0      1
## 11765  7290

accuracy_multi <- 1-mean(class_multi!=new_test$READ)
cat("Accuracy is for multivariate logistic ", accuracy_multi,"\n")

## Accuracy is for multivariate logistic  0.664445

cat("The AUC score is ",roccurve_multi$auc,"\n")

## The AUC score is  0.7588068
```

```
# random forest
```

```
set.seed(16)
ranger_read3 <- ranger(
  formula = READ ~ .,
  data = new_train[,c(3:11,24)],
  num.trees = 1000,
  mtry = 5,
  sample.fraction = .55,
  min.node.size=7,
  importance = "impurity"
)
```


```
## Growing trees.. Progress: 85%. Estimated remaining time: 5 seconds.
```

```
# ranger_read2 <- ranger(
#   formula = READ ~ .,
#   data = new_train[,c(1,2,3,4,5,6,8,9,11,14,24)],
#   num.trees = 1000,
#   mtry = 5,
#   sample.fraction = .55,
#   min.node.size=7,
#   importance = "impurity"
# )
```

```
rangerpre2 <- predict(ranger_read3,new_test,type = "response")
rangerpre2$predictions %>% range
```

```
## [1] 0.0001826687 0.7181991760
```

```
logpre_rf <- rangerpre2$predictions
roccurve_rf <- roc(READ~logpre_rf,data = new_test)
plot(roccurve_rf,xlim=c(0,1),col="blue",asp = NA)
```



```
coords(roc=roccurve_rf,"best","threshold")

## threshold specificity sensitivity
## 0.08616275 0.65980159 0.76677668

classifier_rf <- as.numeric(coords(roccurve_rf,"best","threshold"))
class_rf <- ifelse(logpre_rf>classifier_rf[1],1,0)
table(class_rf)

## class_rf
##      0      1
## 11797  7258

accuracy_rf <- 1-mean(class_rf!=new_test$READ)
cat("Accuracy is for random forest ", accuracy_rf,"\n")

## Accuracy is for random forest  0.6700079

cat("The AUC score is ",roccurve_rf$auc,"\n")

## The AUC score is  0.7808986

# names(new_train[,c(1,2,3,4,5,6,7,8,9,10,11,12,13,14,24)])
new_train %>% names()
```

```

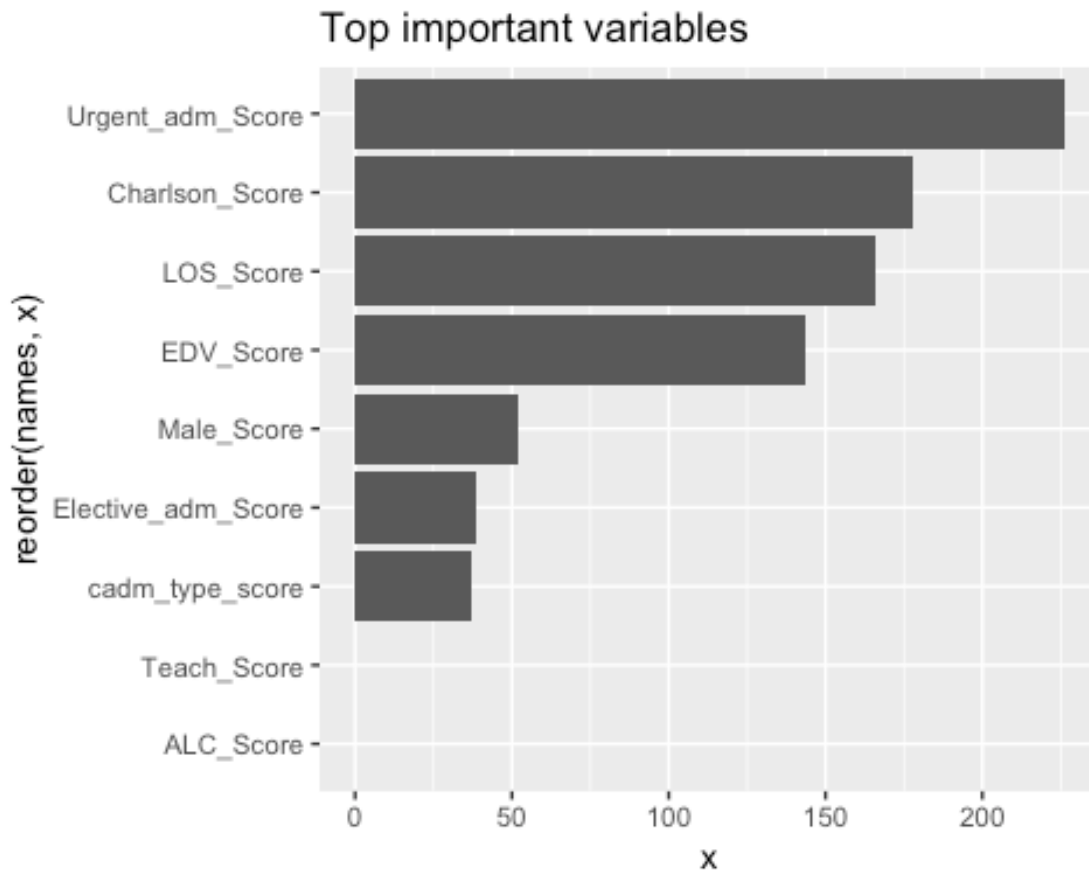
## [1] "Gender" "Age"
## [3] "EDV_Score" "cadm_type_score"
## [5] "LOS_Score" "Charlson_Score"
## [7] "ALC_Score" "Elective_adm_Score"
## [9] "Urgent_adm_Score" "Teach_Score"
## [11] "Male_Score" "Admit_DT"
## [13] "Discharge_DT" "los_days"
## [15] "ADMIT_SOURCE" "dis_from_base_class"
## [17] "loc_name" "DISCH_DISPOSITION"
## [19] "marital_status" "PAT_HOMELESS_YN"
## [21] "zip" "CMS_Readmission_unplanned_flag"
## [23] "Clarity.LACE..Score" "READ"

ranger_read3$variable.importance %>%
  tidy() %>%
  dplyr::arrange(desc(x)) %>%
  dplyr::top_n(25) %>%
  ggplot(aes(reorder(names, x), x)) +
  geom_col() +
  coord_flip() +
  ggtitle("Top important variables")

## Warning: 'tidy.numeric' is deprecated.
## See help("Deprecated")

## Selecting by x

```



```
library(kernlab, warn.conflicts = F)      # SVM methodology
library(e1071)                          # SVM methodology
library(ISLR)                           # contains example data set "Khan"
# install.packages("ISLR")
# install.packages("RColorBrewer")
library(RColorBrewer) # customized coloring of plots

# # sample(1:nrow(new_train), 2000)
# # new_train[sample(new_train, , replace = F), c(1, 2, 3, 4, 5, 6, 8, 9, 11, 14, 24)]
# # svm
# set.seed(16)
# svmfit <- svm(READ~., data = new_train[sample(1:nrow(new_train), 3000
# 0), c(1, 2, 3, 4, 5, 6, 8, 9, 11, 14, 24)], kernel = "linear", scale = FALSE)
# svmpre <- predict(svmfit, new_test)
# svmpre %>% range()
# roccurve_svm <- roc(READ~svmpre, data = new_test)
# plot(roccurve_svm)
# coords(roc=roccurve_svm, "best", "threshold")
# classifier_svm <- as.numeric(coords(roccurve_svm, "best", "threshold"))
# class_svm <- ifelse(svmpre > classifier_svm[1], 1, 0)
# table(class_svm)
# accuracy_svm <- 1 - mean(class_svm != new_test$READ)
```

```

# cat("Accuracy is for SVM ", accuracy_svm, "\n")
# cat("The AUC score is ", roccurve_svm$auc, "\n")
# ```
#
# ```{r}
#
# set.seed(16)
# svmfit4 <- svm(READ~., data = new_train[sample(1:nrow(new_train)),c(1,
2,3,4,5,6,8,9,11,14,24)], kernel = "radial", gamma = 1, cost = 1)
# svmpre <- predict(svmfit4,new_test)
# svmpre %>% range()
# roccurve_svm <- roc(READ~svmpre,data = new_test)
# plot(roccurve_svm,xlim=c(0,1),col="blue",asp = NA)
# coords(roc=roccurve_svm , "best", "threshold")
# classifier_svm <- as.numeric(coords(roccurve_svm, "best", "threshold"))
# class_svm <- ifelse(svmpre>classifier_svm[1],1,0)
# table(class_svm)
# # table(new_test$READ)
# accuracy_svm <- 1-mean(class_svm!=new_test$READ)
# cat("The Accuracy is for SVM ", accuracy_svm, "\n")
# cat("The AUC score is ", roccurve_svm$auc, "\n")
#
# misclass <- table(predict = class_svm, truth = new_test$READ)

library(rsample)      # data splitting
library(dplyr)        # data wrangling
library(rpart)        # performing regression trees
library(rpart.plot)   # plotting regression trees
# install.packages("rpart.plot") # plotting regression trees
library(ipred)        # bagging
library(caret)        # bagging

# bagging decision trees.

# set.seed(16)
# bagged_m1 <- bagging(
#   formula = READ~.,
#   data    = new_train[sample(1:nrow(new_train)),c(1,2,3,4,5,6,8,9,11,
14,24)],
#   coob    = TRUE,
#   nbagg   = 31
# )
# bagpre <- predict(bagged_m1,new_test)
# bagpre %>% range()
# roccurve_bag <- roc(READ~bagpre,data = new_test)
# plot(roccurve_bag,xlim=c(0,1),col="blue",asp = NA)
# coords(roc=roccurve_bag , "best", "threshold")
# classifier_bag <- as.numeric(coords(roccurve_bag, "best", "threshold"))
# class_bag <- ifelse(bagpre>classifier_bag[1],1,0)
# table(class_bag)

```

```

# table(new_test$READ)
# accuracy_bag <- 1-mean(class_bag!=new_test$READ)
# cat("Accuracy is for Bootstrap aggregating (bagging) ", accuracy_bag,
"\n")
# cat("The AUC score is ",roccurve_bag$auc,"\n")
# plot(varImp(bagged_m1), 20)

# set.seed(16)
# ctrl <- trainControl(method = "cv", number = 3)
# # bagged_cv <- train(
# #   factor(READ) ~ .,
# #   data = new_train[sample(1:nrow(new_train)),c(1,2,3,4,5,6,8,9,11,1
4,24)],
# #   method = "treebag",
# #   trControl = ctrl,
# #   importance = TRUE
# # )
# set.seed(16)
# bagged_cv2 <- train(
#   factor(READ) ~ .,
#   data = new_train[sample(1:nrow(new_train),2000),c(3:11,24)],
#   method = "treebag",
#   trControl = ctrl,
#   importance = TRUE
# )
# bagpre <- predict(bagged_cv,new_test,type="prob")
# bagpre
# table(as.integer(bagpre)-1)
# l=(as.integer(bagpre)-1)
# roccurve_bag <- roc(READ~ bagpre[,2],data = new_test)
# plot(roccurve_bag,xlim=c(0,1),col="blue",asp = NA)
# auc(roccurve_bag)
# coords(roc=roccurve_bag ,"best","threshold")
# classifier_bag <- as.numeric(coords(roccurve_bag,"best","threshold"))
# class_bag <- ifelse(bagpre>classifier_bag[1],1,0)
# table(class_bag)
# table(new_test$READ)
# accuracy_bag <- 1-mean(bagpre!=new_test$READ)
# cat("Accuracy is for Bootstrap aggregating (bagging) ", accuracy_bag,
"\n")
# cat("The AUC score is ",roccurve_bag$auc,"\n")
# plot(varImp(bagged_m1), 20)
#
#
#
# plot(varImp(bagged_cv2),10)
#
# test.bagprob=bagpre
# library(ipred)

```

```

# #prepare bagged model for curve
# test.bagprob = predict(train_bag, type = "prob", newdata = test)
# bagpred = prediction( bagpre[,2], test$class)
# bagperf = performance(bagpre,"tpr","fpr")
# bagprek <- ifelse(bagpre[,2]>bagpre[,1],bagpre[,2],bagpre[,1])
# > plot(perf, main="ROC", colorize=T)
# > plot(bagperf, col=2, add=TRUE)
# > plot(perf, col=1, add=TRUE)
# > legend(0.6, 0.6, c('ctree', 'bagging'), 1:2)
# ```
# ```{r}
# # # assess 10-50 bagged trees
# # ntree <- 25:40
# #
# # # create empty vector to store OOB RMSE values
# # rmse <- vector(mode = "numeric", length = length(ntree))
# #
# # for (i in seq_along(ntree)) {
# #   # reproducibility
# #   set.seed(123)
# #
# #   # perform bagged model
# #   model <- bagging(
# #     formula = READ~.,
# #     data     = new_train[sample(1:nrow(new_train)),c(1,2,3,4,5,6,8,9,1
# 1,14,24)],
# #     coob     = TRUE,
# #     nbagg    = ntree[i]
# #   )
# #   # get OOB error
# #   rmse[i] <- model$err
# # }
# # plot(ntree, rmse, type = 'l', lwd = 2)
# # abline(v = 25, col = "red", lty = "dashed")
# # #
#
# 1,2,3,4,5,6,7,8,9,10,11,24
###numeric only
features_train <- as.matrix(new_train[,c(3:11)])
response_train <- as.matrix(new_train[,24])
# names(new)

features_test <- as.matrix(new_test[,c(3:11)])
response_test <- as.matrix(new_test[,24])

library("xgboost")

##
## Attaching package: 'xgboost'

```



```

## The following object is masked from 'package:dplyr':
##
##      slice


# GBM xgboost: Training and tuning with the xgboost package
params <- list(
  eta = 0.1,
  max_depth = 9,
  min_child_weight = 3,
  subsample = 1,
  colsample_bytree = 0.8
)
set.seed(16)
xgb.fit.final <- xgboost(
  params = params,
  data = features_train,
  label = response_train,
  nrounds = 46,
  objective = "reg:linear",
  verbose = 0
)

gbmpre <- predict(xgb.fit.final, features_test)
gbmpre %>% range()

## [1] -0.03919309  0.90631318

roccurve_gbm <- roc(new_test$READ~gbmpre)
plot(roccurve_gbm, xlim=c(0,1), col="blue", asp = NA)

```



```
coords(roc=roccurve_gbm , "best", "threshold")

##   threshold specificity sensitivity
## 0.09373079 0.65278181 0.79262926

classifier_gbm <- as.numeric(coords(roccurve_gbm, "best", "threshold"))
class_gbm <- ifelse(gbm$pre > classifier_gbm[1], 1, 0)
table(class_gbm)

## class_gbm
##      0      1
## 11629  7426

# table(response_test)s
accuracy_bgm <- 1 - mean(class_gbm != new_test$READ)
cat("Accuracy is for Gradient Boosting ", accuracy_bgm, "\n")


## Accuracy is for Gradient Boosting 0.6661244

cat("The AUC score is ", roccurve_gbm$auc, "\n")

## The AUC score is 0.7907732

# create importance matrix
importance_matrix <- xgb.importance(model = xgb.fit.final)
```

```
# variable importance plot
xgb.plot.importance(importance_matrix, top_n = 13, measure = "Gain")
```



```
## cat("The Accuracy is for simple logistic ", accuracy,"\n")
## cat("The AUC score is ",roccurve$auc,"\n")
# cat("The Accuracy is for multivariate logistic ", accuracy_multi,"\n")
# cat("The AUC score is ",roccurve_multi$auc,"\n")
# cat("The Accuracy is for random forest ", accuracy_rf,"\n")
# cat("The AUC score is ",roccurve_rf$auc,"\n")
## cat("The Accuracy is for SVM ", accuracy_svm,"\n")
## cat("The AUC score is ",roccurve_svm$auc,"\n")
# cat("The Accuracy is for Bagging ", accuracy_bag,"\n")
# cat("The AUC score is ",roccurve_bag$auc,"\n")
# cat("The Accuracy is for Gradient Boosting ", accuracy_bgm,"\n")
# cat("The AUC score is ",roccurve_gbm$auc,"\n")
## plot(ranger_read$variable.importance)
```

Balance data

```
library(DMwR,warn.conflicts = F)
```

```
## Loading required package: grid
```

```

new_train$READ <- as.factor(new_train$READ)
new_train <- SMOTE(READ~.,new_train[,c(1,2,3,4,5,6,8,9,11,14,24)],perc.
over = 100,perc.under = 200)
new_train$READ <- as.numeric(new_train$READ)

head(new_train,2)
## 76219                6                30                0                70                1
## 31487.1              0                18                0                 2                1
## [ reached 'max' / getOption("max.print") -- omitted 12962 rows ]

# library(xgboost)
#
# # grid search
# for(i in 1:nrow(hyper_grid)) {
#
#   # create parameter list
#   params <- list(
#     eta = hyper_grid$eta[i],
#     max_depth = hyper_grid$max_depth[i],
#     min_child_weight = hyper_grid$min_child_weight[i],
#     subsample = hyper_grid$subsample[i],
#     colsample_bytree = hyper_grid$colsample_bytree[i]
#   )
#
#   # reproducibility
#   set.seed(123)
#
#   # train model
#   xgb.tune <- xgb.cv(
#     params = params,
#     data = features_train,
#     label = response_train,
#     nrounds = 5000,
#     nfold = 5,
#     objective = "reg:linear", # for regression models
#     verbose = 0,              # silent,
#     early_stopping_rounds = 10 # stop if no improvement for 10 consec
#   )
#
#   # add min training error and trees to grid
#   hyper_grid$optimal_trees[i] <- which.min(xgb.tune$evaluation_log$te
#   st_rmse_mean)
#   hyper_grid$min_RMSE[i] <- min(xgb.tune$evaluation_log$test_rmse_mea
#   n)
# }
#
# hyper_grid %>%
#   dplyr::arrange(min_RMSE) %>%
#   head(10)

```

```

# create hyperparameter grid
hyper_grid <- expand.grid(
  eta = 0.1,
  max_depth = 7,
  min_child_weight = 3,
  subsample = 1,
  colsample_bytree = 0.8,
  optimal_trees = 0,          # a place to dump results
  min_RMSE = 0               # a place to dump results
)

# hyper_grid <- expand.grid(
#   eta = c(.01, .05, .1, .3),
#   max_depth = c(1, 3, 5, 7),
#   min_child_weight = c(1, 3, 5, 7),
#   subsample = c(.65, .8, 1),
#   colsample_bytree = c(.8, .9, 1),
#   optimal_trees = 0,          # a place to dump results
#   min_RMSE = 0               # a place to dump results
# )

# grid search
for(i in 1:nrow(hyper_grid)) {
  #
  #   # create parameter list
  #   params <- list(
  #     eta = hyper_grid$eta[i],
  #     max_depth = hyper_grid$max_depth[i],
  #     min_child_weight = hyper_grid$min_child_weight[i],
  #     subsample = hyper_grid$subsample[i],
  #     colsample_bytree = hyper_grid$colsample_bytree[i]
  #   )
  #
  #   # reproducibility
  #   set.seed(123)
  #
  #   # train model
  #   xgb.tune <- xgb.cv(
  #     params = params,
  #     data = features_train,
  #     label = response_train,
  #     nrounds = 5000,
  #     nfold = 5,
  #     objective = "reg:linear", # for regression models
  #     verbose = 0,              # silent,
  #     early_stopping_rounds = 10 # stop if no improvement for 10 consec
  #   )
  #
  #   # add min training error and trees to grid

```

```
# hyper_grid$optimal_trees[i] <- which.min(xgb.tune$evaluation_log$test_rmse_mean)
# hyper_grid$min_RMSE[i] <- min(xgb.tune$evaluation_log$test_rmse_mean)
# }
#
# hyper_grid %>%
#   dplyr::arrange(min_RMSE) %>%
#   head(10)
#46 54 55
```